

Overview of Optimization for ML

Sunita Sarawagi

Motivation

- Many ML training algorithms can be posed as a continuous optimization problem

$$\min_{\substack{\{w_1, \dots, w_d\} \\ w_j \in \mathbb{R}}} F(w_1, w_2, \dots, w_d) \Leftrightarrow \min_{\substack{w \\ w \in \mathbb{R}^d}} F(\vec{w})$$

where the variables are the parameters of the model

$F(w)$: loss over training data + regularizer

- Loss = continuous approximation of 0/1 error.
- Finding the minima of general functions could be intractably hard
 - Doable for certain types of functions \rightarrow convex functions.

Examples of functions

$d=1$

$$F(w_1) = 5w_1 + 7 \quad ; \quad F(w_1) = \log(1 + e^{-w_1})$$



$$F(w_1) = \sum_{i=1}^N (y_i - w_1 x^i)^2 \quad \leftarrow$$

$(x^i, y_i) \equiv$ known constant

$d=2$

$$F(w_1, w_2) = w_1^2 + 2w_2 + w_1 w_2 + w_2^3 + \log(1 + w_1)$$

Convex functions




A function $F(w)$ is convex in w if and only if (iff) for any $w^1, w^2 \in \mathbb{R}^d, \lambda \in [0,1]$, $F(\lambda w^1 + (1 - \lambda)w^2) \leq \lambda F(w^1) + (1 - \lambda)F(w^2)$,
convex combination


Concave functions = negative of convex functions.

Convex Vs Non-convex functions in 2-d


Convex Function: $f(w_1, w_2) = w_1^2 + w_2^2$



$f(w_1, w_2) = w_1^3 - 3w_1w_2^2$




$f(w_1, w_2) = 20 + w_1^2 + w_2^2 - 10 * (\cos(2\pi w_1) + \cos(2\pi w_2))$

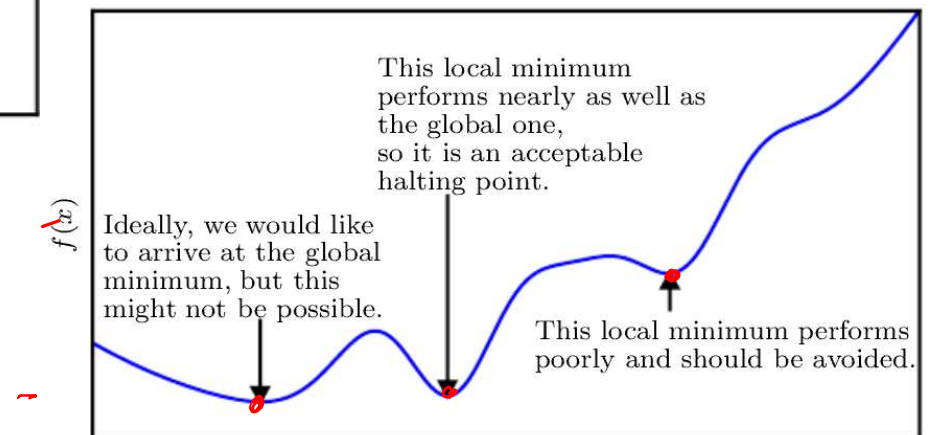


Optimizing a function


- $\min_w F(w), w \in R^d$
- A point w^* is a **global minima** of $F(w)$ if $F(w) \geq F(w^*)$
- A point w^* is a **local minima** of $F(w)$ if there exists an $\epsilon > 0$ such that $F(w) \geq F(w^*) \quad \forall |w - w^*| \leq \epsilon$






$d = 1$



Local and global minima in 2-D





Gradient of a differentiable function

- Derivative of a function on a single variable measures the rate of change of the function.

$$F(w_1) = w_1^2 + 2w_1^3$$

$$F(w_1) = \log(1 + e^{-w_1})$$

$$\frac{\partial F(w_1)}{\partial w_1} = 2w_1 + 6w_1^2$$

$$\frac{\partial F(w_1)}{\partial w_1} = \frac{-e^{-w_1}}{1 + e^{-w_1}}$$

$$F(w_1) = \sum_{i=1}^N (y^i - w_1 x^i)^2$$

$$\frac{\partial F(w_1)}{\partial w_1} = \sum_{i=1}^N 2(y^i - w_1 x^i)(-x^i)$$

Gradients of multi-variable functions

- For multivariable functions $F(w_1, \dots, w_d)$ we can define partial derivative of F w.r.t each of the variables.

$$\frac{\partial F}{\partial w_1}, \frac{\partial F}{\partial w_2}, \dots, \frac{\partial F}{\partial w_d}$$

- Gradient of $F(w)$ denoted as $\nabla F(w)$: vector of partial derivative of the function

$$\nabla_w F(\vec{w}) = \begin{bmatrix} \frac{\partial F}{\partial w_1} \\ \frac{\partial F}{\partial w_2} \\ \vdots \\ \frac{\partial F}{\partial w_d} \end{bmatrix}$$

$d=2$ Example of gradient

$$F(w_1, w_2) = \underline{F(\mathbf{w})} = \frac{1}{2} (w_1^2 + 10 w_2^2) \leftarrow$$


$$\nabla F(\mathbf{w}) = \begin{bmatrix} \frac{\partial F}{\partial w_1} \\ \frac{\partial F}{\partial w_2} \end{bmatrix} = \begin{bmatrix} \frac{1}{2} \cdot 2 \cdot w_1 \\ \frac{1}{2} \cdot 10 \cdot 2 \cdot w_2 \end{bmatrix} = \begin{bmatrix} w_1 \\ 10 w_2 \end{bmatrix}$$

$$\nabla F(\mathbf{w}) \big|_{\mathbf{w} = \underline{[1, 3]'}} = \begin{bmatrix} 1 \\ 10 \times 3 \end{bmatrix}.$$

Minima and gradients

- **Theorem:** If $F(w)$ is differentiable and if w^* is a local minima of $F(w)$, then $\nabla F(w^*) = 0$
- Theorem: For convex functions a w^* is a global minima if and only if $\nabla F(w^*) = 0$. That is the local minima is the global minima.


Visualizing gradients in 2-D



$$f(w_1) = (w_1 - 1)^2$$
$$\left. \frac{\partial f(w_1)}{\partial w_1} \right|_{w_1=2} = 2(w_1 - 1)$$

$$= 2(2 - 1)$$

$$\left. \frac{\partial f(w_1)}{\partial w_1} \right|_{w_1 = -1} = 2$$
$$= 2(-1 - 1)$$
$$= -4$$



Iterative optimization algorithms

- Most often we will not be able to solve for $\nabla F(w) = 0$ in closed form
 - E.g. MLE for the logistic loss. (Show)
- Iterative algorithms (General template)
 - w^0 = Choose an initial point.
 - For $t = 1$ to stopping criteria (local minima, maximum iterations, etc)
 - $w^{t+1} \leftarrow$ Move to a near-by point w^t such that $F(w)$ reduces.
- Many iterative algorithms have been proposed for such cases
 - Zero-th order algorithm (line-search) for optimizing 1-d convex functions.
 - First-order or gradient-based algorithms
 - Gradient descent
 - Conjugate gradient descent,
 - .
 - Second-order algorithms
 - Semi or pseudo second order algorithms

Example of solving in closed form

$$F(w_1, w_2) = (w_1 - 1)^2 + 3(w_2 - 10)^2$$

$$\nabla F(w) = \begin{bmatrix} 2(w_1 - 1) \\ 6(w_2 - 10) \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$\begin{aligned} \Rightarrow 2w_1^* - 2 &= 0 & \Rightarrow w_1^* &= 1 \\ 6w_2^* - 60 &= 0 & \Rightarrow w_2^* &= 10 \end{aligned}$$

$$\begin{aligned} F(w_1, w_2) &= \log(e^{-w_1 + 2w_2}) \\ &+ \log(e^{w_1 + 7w_2}) \end{aligned}$$

Iterative Optimization using gradients

- Choose an arbitrary initial point : $w^0 \leftarrow [0 \dots 0]^T$
- λ = Chosen learning rate \leftarrow has to be small but not too small.
- Epoch $t = 0$
- While stopping criteria not reached (t)

$\nabla F(w) \big|_{w=w^t}$ $\nabla F(w^t) \equiv$ Compute gradient of function at current w^t .

if $\|\nabla F(w^t)\| \approx 0$, then return w^t as minima.

$$w^{t+1} = w^t - \lambda \nabla F(w^t)$$

\nwarrow
descend gradient


$$t = t + 1$$

$$F(w^{t+1}) \leq F(w^t)$$

if $F(w)$ is differentiable
and λ is small.

Iterative minimization of 1-d convex functions

- Convex functions \Rightarrow double derivative (rate of change of gradients) is non-negative. Minima where derivation = 0.



Example: gradient descent

- $F(w_1, w_2) = \frac{1}{2}(w_1^2 + 10w_2^2)$

- $w^0 = [10, 1]^T$

$$\nabla F(w^0) = [10, 10]^T; \quad \lambda = 0.1 \quad F(w^0) = \frac{110}{2}$$

$$w^1 = \begin{bmatrix} 10 \\ 1 \end{bmatrix} - 0.1 \begin{bmatrix} 10 \\ 10 \end{bmatrix} = \begin{bmatrix} 9 \\ 0 \end{bmatrix} \quad F(w^1) = \frac{81}{2}$$

$$\nabla F(w^1) = \begin{bmatrix} 9 \\ 0 \end{bmatrix}$$

$$w^2 = \begin{bmatrix} 9 \\ 0 \end{bmatrix} - 0.1 \begin{bmatrix} 9 \\ 0 \end{bmatrix} = \begin{bmatrix} 8.1 \\ 0 \end{bmatrix}$$

\vdots

$$w^T = \begin{bmatrix} 0.01 \\ 0 \end{bmatrix}$$

$$\nabla F(w^T) = \left\| \begin{bmatrix} 0.01 \\ 0 \end{bmatrix} \right\| \approx 0 \rightarrow \text{stop}$$

Gradient descent geometrically

- See [this colab link](#) to run the demo yourself.

Stochastic gradient descent

- Stochastic approximation to gradient descent optimization when applied over sum of errors on several i.i.d training examples
- Typical training objective:
 - $L(w) = \frac{1}{N} \sum_{i=1} L(f(x^i; w), y_i)$
 - True gradient:
 - Stochastic approximation:
- More efficient than full batch
- Empirical found to be better at optimizing non-convex functions because of noisy nature of gradients.